

ANN in Deep Learning (MNIST with TensorFlow/Keras)

Muhammad Zaid Kamil

Installation

```
In [5]: import tensorflow.compat.v1 as tf

In [6]: tf.disable_v2_behavior()

WARNING:tensorflow:From C:\Users\zombi\AppData\Roaming\Python\Python39\site-packages\tensorflow\python\compat\v2_compat.py:107:
disable_resource_variables (from tensorflow.python.ops.variable_scope) is deprecated and will be removed in a future version.
Instructions for updating:
non-resource variables are not supported in the long term

In [7]: import tensorflow as tf

In [8]: hello = tf.constant('Hello, TensorFlow!')

In [9]: sess = tf.compat.v1.Session()

In [10]: print(sess.run(hello))

b'Hello, TensorFlow!'
```

MNIST with TensorFlow

Read the description of the task in <https://www.tensorflow.org/tutorials/quickstart/beginner>.

Step 1:

```
In [2]: import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten(input_shape=(28,28)))
model.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
model.add(tf.keras.layers.Dropout(0.2))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

In [3]: predictions = model(x_train[:1]).numpy()
predictions

Out[3]: array([[0.09971632, 0.10760017, 0.16850476, 0.13726926, 0.05614685,
0.11357612, 0.0781882 , 0.05929397, 0.11647072, 0.0632336 ]],
dtype=float32)
```

Step 1: Setting up Tensor Flow

Step 2: Loading the dataset

Step 3: Building a ML model

Step 5: Predictions to return logits one for each class

```
In [4]: tf.nn.softmax(predictions).numpy()
```

```
Out[4]: array([[0.09991138, 0.10070218, 0.10702603, 0.10373469, 0.09565176,
               0.10130578, 0.09778346, 0.09595326, 0.10159943, 0.09633204]],
          dtype=float32)
```

Step 6: The `tf.nn.softmax` is used to convert logits → probabilities for each class

```
In [5]: loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
```

Step 7: Define Loss function

```
In [6]: loss_fn(y_train[:1], predictions).numpy()
```

```
Out[6]: 2.2896118
```

Step 8: Untrained model, initial loss close to 2.8

```
In [8]: model.fit(x_train, y_train, epochs=5)
```

```
Epoch 1/5
1875/1875 [=====] - 5s 2ms/step - loss: 1.5796 - accuracy: 0.9007
Epoch 2/5
1875/1875 [=====] - 4s 2ms/step - loss: 1.5224 - accuracy: 0.9446
Epoch 3/5
1875/1875 [=====] - 4s 2ms/step - loss: 1.5090 - accuracy: 0.9565
Epoch 4/5
1875/1875 [=====] - 4s 2ms/step - loss: 1.5020 - accuracy: 0.9625
Epoch 5/5
1875/1875 [=====] - 4s 2ms/step - loss: 1.4971 - accuracy: 0.9668
```

```
Out[8]: <keras.callbacks.History at 0x1b7c32989a0>
```

Step 9: Train and evaluate the model. `Model.fit` to adjust model parameters and adjust loss

```
In [9]: model.evaluate(x_test, y_test, verbose=2)
```

```
313/313 - 1s - loss: 1.4923 - accuracy: 0.9717 - 654ms/epoch - 2ms/step
```

```
Out[9]: [1.492332577705383, 0.9717000126838684]
```

Step 9: The `Model.evaluate` method checks the models performance, usually on a "Validation-set" or "Test-set".

Step 10: The image classifier is trained to 97% accuracy in this dataset

```
In [10]: probability_model = tf.keras.Sequential([
        model,
        tf.keras.layers.Softmax()
    ])
```

Step 11: Wrap the trained model and attach to softmax

```
In [11]: probability_model(x_test[:5])
```

```
Out[11]: <tf.Tensor: shape=(5, 10), dtype=float32, numpy=
array([[0.08533674, 0.08533674, 0.08533674, 0.08533674, 0.08533674,
        0.08533674, 0.08533674, 0.23196931, 0.08533674, 0.08533674],
       [0.08533674, 0.08533674, 0.23196931, 0.08533674, 0.08533674,
        0.08533674, 0.08533674, 0.08533674, 0.08533674, 0.08533674],
       [0.0853368 , 0.23196824, 0.08533701, 0.0853368 , 0.0853368 ,
        0.0853368 , 0.08533687, 0.08533693, 0.08533686, 0.0853368 ],
       [0.23196931, 0.08533674, 0.08533674, 0.08533674, 0.08533674,
        0.08533674, 0.08533674, 0.08533674, 0.08533674, 0.08533674],
       [0.08539622, 0.08539622, 0.08539622, 0.08539622, 0.23102519,
        0.08539622, 0.08539622, 0.08539643, 0.08539622, 0.08580474]],
        dtype=float32)>
```

Step 12: The Model to return probability

Create a function that is a model for recognizing digits, based on looking at every pixel in the image

```
In [12]: import tensorflow as tf
         tf.__version__
```

```
Out[12]: '2.9.1'
```

```
In [13]: mnist = tf.keras.datasets.mnist
```

```
In [14]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Step 1: Retrieving Data set of 28x28 images of handwritten digits of 0-9. Image.


#. Feed through the pictures in neural network and neural network check which number image

```
In [15]: import matplotlib.pyplot as plt
          print(x_train[0])
          plt.show()
```

[illegible]


Step 2: The number plot shown represents pixels range from 0 to 255 pixels. Pixels > 0 represents the digit

```
In [17]: import matplotlib.pyplot as plt
plt.imshow(x_train[0])
plt.show()
```



```
In [18]: plt.imshow(x_train[0], cmap = plt.cm.binary)
```


```
Out[18]: <matplotlib.image.AxesImage at 0x236460b4f10>
```



Step 3: The actual digit representation image in Black and White

```
In [21]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = tf.keras.utils.normalize(x_train, axis = 1)
x_test = tf.keras.utils.normalize(x_test, axis = 1)
```

```
In [22]: plt.imshow(x_train[0], cmap = plt.cm.binary)
plt.show()
print(x_train[0])
```



```
[0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0.00393124 0.02332955 0.02620568 0.02625207 0.17420356 0.17566281
 0.28629534 0.05664824 0.51877786 0.71632322 0.77892406 0.89301644
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0.05780486 0.06524513 0.16128198 0.22713296
 0.22277047 0.32790981 0.36833534 0.3689874 0.34978968 0.32678448
 0.368094 0.3747499 0.79066747 0.67980478 0.61494005 0.45002403
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0.12250613 0.45858525 0.45852825 0.43408872 0.37314701
 0.33153488 0.32790981 0.36833534 0.3689874 0.34978968 0.32420121
 0.15214552 0.17865984 0.25626376 0.1573102 0.12298801 0.]
```

Step 4: Normalizing the data. The pixels range from 0 to 1.


```
In [1]: import tensorflow as tf

mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = tf.keras.utils.normalize(x_train, axis = 1)
x_test = tf.keras.utils.normalize(x_test, axis = 1)

model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(128, activation=tf.nn.relu))
model.add(tf.keras.layers.Dense(10, activation=tf.nn.softmax))

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
model.fit(x_train, y_train, epochs=3)

Epoch 1/3
1875/1875 [=====] - 8s 3ms/step - loss: 0.2613 - accuracy: 0.9227
Epoch 2/3
1875/1875 [=====] - 6s 3ms/step - loss: 0.1103 - accuracy: 0.9654
Epoch 3/3
1875/1875 [=====] - 7s 4ms/step - loss: 0.0755 - accuracy: 0.9764

Out[1]: <keras.callbacks.History at 0x20722261d90>
```

Average of 96% Accuracy of the Neural network after 3 EPOCHS. This was in a sample. – model generalized

```
In [2]: val_loss, val_acc = model.evaluate(x_test, y_test)
print(val_loss, val_acc)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.1000 - accuracy: 0.9678
0.10002856701612473 0.9678000211715698
```

Calculating validation Loss = 0.1 and accuracy = 0.9678

```
In [3]: predictions = model.predict([x_test])
```

```
WARNING:tensorflow:Layers in a Sequential model should only have a single input tensor. Received: inputs=(<tf.Tensor 'IteratorGetNext:0' shape=(None, 28, 28) dtype=float32>,. Consider rewriting this model with the Functional API.
313/313 [=====] - 1s 3ms/step
```

```
In [4]: print(predictions)
```

```
[[3.0582585e-09 3.4056779e-07 1.1552959e-06 ... 9.9999136e-01
 3.9908922e-08 1.8685662e-06]
 [5.1894166e-11 3.5538431e-03 9.9644476e-01 ... 1.4328446e-09
 1.3937509e-07 4.0627085e-10]
 [8.1326927e-09 9.9993402e-01 1.7088509e-05 ... 1.3252626e-05
 3.1097064e-05 9.1889092e-08]
 ...
 [1.7900863e-10 2.9408066e-06 5.2566147e-09 ... 1.5657708e-05
 4.7391813e-06 2.0817497e-04]
 [1.7899889e-05 2.9049856e-06 4.4617636e-07 ... 1.9866290e-07
 2.0838568e-03 9.6705639e-07]
 [6.4102355e-09 2.7654480e-07 1.0818975e-07 ... 8.1363150e-13
 6.4531747e-08 2.4956769e-08]]
```

```
In [6]: import numpy
        print(numpy.argmax(predictions[0]))
```

7

Classify Images of Clothing

```
In [1]: # TensorFlow and tf.keras
import tensorflow as tf

# Helper Libraries
import numpy as np
import matplotlib.pyplot as plt

print(tf.__version__)
```

2.10.0

```
In [2]: fashion_mnist = tf.keras.datasets.fashion_mnist

        (train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data()
```

```
In [3]: class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
                        'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
```

```
In [4]: train_images.shape
```

Out[4]: (60000, 28, 28)

```
In [5]: len(train_labels)
```

Out[5]: 60000

```
In [6]: train_labels
```

Out[6]: array([9, 0, 0, ..., 3, 0, 5], dtype=uint8)


```
In [7]: test_images.shape
```

Out[7]: (10000, 28, 28)

```
In [8]: len(test_labels)
```

Out[8]: 10000


```
In [9]: plt.figure()
plt.imshow(train_images[0])
plt.colorbar()
plt.grid(False)
plt.show()
```



```
In [10]: train_images = train_images / 255.0
test_images = test_images / 255.0
```

```
In [11]: plt.figure(figsize=(10,10))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_images[i], cmap=plt.cm.binary)
    plt.xlabel(class_names[train_labels[i]])
plt.show()
```



Ankle boot



T-shirt/top



T-shirt/top



Dress



T-shirt/top



```
In [12]: model = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10)
])
```

```
In [13]: model.compile(optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'])
```

```
In [13]: model.compile(optimizer='adam',
                    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                    metrics=['accuracy'])
```

```
In [14]: model.fit(train_images, train_labels, epochs=10)

Epoch 1/10
1875/1875 [=====] - 4s 2ms/step - loss: 0.5034 - accuracy: 0.8244
Epoch 2/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.3828 - accuracy: 0.8610
Epoch 3/10
1875/1875 [=====] - 3s 1ms/step - loss: 0.3400 - accuracy: 0.8773
Epoch 4/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.3157 - accuracy: 0.8843
Epoch 5/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.2961 - accuracy: 0.8898
Epoch 6/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.2795 - accuracy: 0.8961
Epoch 7/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.2694 - accuracy: 0.8996
Epoch 8/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.2588 - accuracy: 0.9028
Epoch 9/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.2498 - accuracy: 0.9067
Epoch 10/10
1875/1875 [=====] - 3s 2ms/step - loss: 0.2412 - accuracy: 0.9097
```

```
Out[14]: <keras.callbacks.History at 0x2221ae7e0d0>
```

```
In [15]: test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)

print('\nTest accuracy:', test_acc)

313/313 - 1s - loss: 0.3257 - accuracy: 0.8885 - 501ms/epoch - 2ms/step

Test accuracy: 0.8884999752044678
```

```
In [16]: probability_model = tf.keras.Sequential([model,
                    tf.keras.layers.Softmax()])
```

```
In [17]: predictions = probability_model.predict(test_images)

313/313 [=====] - 0s 1ms/step
```

```
In [18]: predictions[0]
```

```
Out[18]: array([1.80547886e-05, 1.80332869e-08, 1.09664903e-07, 2.56643684e-08,
                1.06003222e-06, 5.15667023e-04, 3.11436634e-06, 1.26864975e-02,
                1.07358716e-07, 9.86775398e-01], dtype=float32)
```

```
In [19]: np.argmax(predictions[0])
```

```
Out[19]: 9
```

```
In [20]: test_labels[0]
```

```
Out[20]: 9
```

```

In [21]: def plot_image(i, predictions_array, true_label, img):
    true_label, img = true_label[i], img[i]
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])

    plt.imshow(img, cmap=plt.cm.binary)

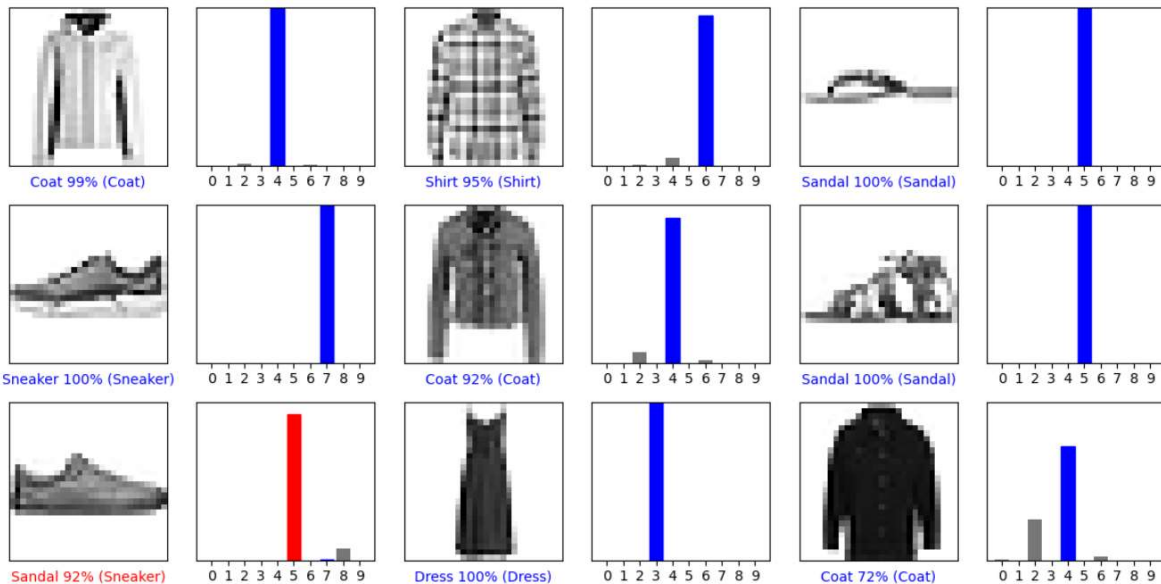
    predicted_label = np.argmax(predictions_array)
    if predicted_label == true_label:
        color = 'blue'
    else:
        color = 'red'

    plt.xlabel("{} {:2.0f}% ({})"
               .format(class_names[predicted_label],
                       100*np.max(predictions_array),
                       class_names[true_label]),
               color=color)


def plot_value_array(i, predictions_array, true_label):
    true_label = true_label[i]
    plt.grid(False)
    plt.xticks(range(10))
    plt.yticks([])
    thisplot = plt.bar(range(10), predictions_array, color="#777777")
    plt.ylim([0, 1])
    predicted_label = np.argmax(predictions_array)

    thisplot[predicted_label].set_color('red')
    thisplot[true_label].set_color('blue')


```



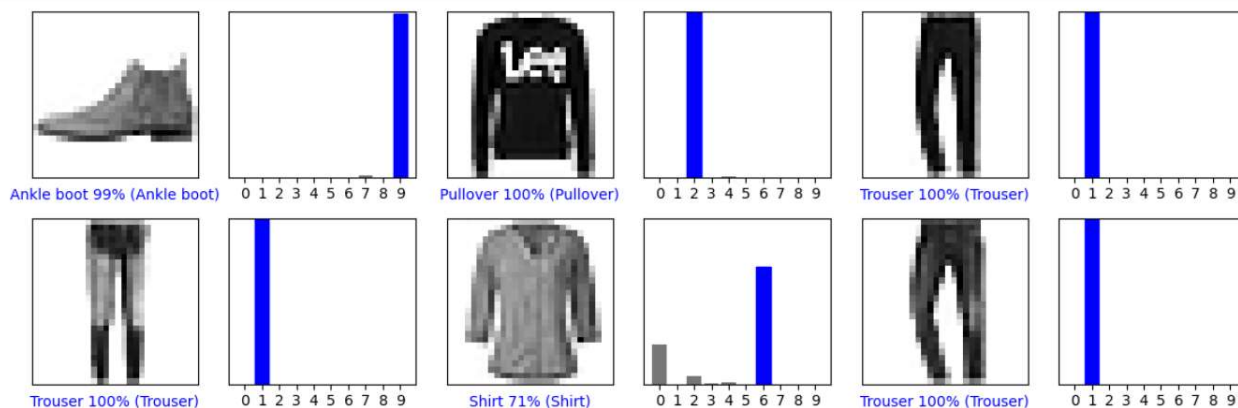
```
In [22]: i = 0
plt.figure(figsize=(6,3))
plt.subplot(1,2,1)
plot_image(i, predictions[i], test_labels, test_images)
plt.subplot(1,2,2)
plot_value_array(i, predictions[i], test_labels)
plt.show()
```



```
In [23]: i = 12
plt.figure(figsize=(6,3))
plt.subplot(1,2,1)
plot_image(i, predictions[i], test_labels, test_images)
plt.subplot(1,2,2)
plot_value_array(i, predictions[i], test_labels)
plt.show()
```



```
In [24]: # Plot the first X test images, their predicted Labels, and the true Labels.
# Color correct predictions in blue and incorrect predictions in red.
num_rows = 5
num_cols = 3
num_images = num_rows*num_cols
plt.figure(figsize=(2*2*num_cols, 2*num_rows))
for i in range(num_images):
    plt.subplot(num_rows, 2*num_cols, 2*i+1)
    plot_image(i, predictions[i], test_labels, test_images)
    plt.subplot(num_rows, 2*num_cols, 2*i+2)
    plot_value_array(i, predictions[i], test_labels)
plt.tight_layout()
plt.show()
```



In [25]: *# Grab an image from the test dataset.*

```
img = test_images[1]
```

```
print(img.shape)
```

(28, 28)

In [26]: *# Add the image to a batch where it's the only member.*

```
img = (np.expand_dims(img,0))
```

```
print(img.shape)
```

(1, 28, 28)

In [27]: *# Add the image to a batch where it's the only member.*

```
img = (np.expand_dims(img,0))
```

```
print(img.shape)
```

(1, 1, 28, 28)


In [28]: predictions_single = probability_model.predict(img)

```
print(predictions_single)
```

1/1 [=====] - 0s 46ms/step

```
[[1.1324369e-04 4.5422674e-13 9.9893767e-01 4.1311107e-11 7.7292504e-04  
 3.1105825e-13 1.7616725e-04 8.3188371e-18 2.0409157e-09 7.8077664e-17]]
```

In [29]: plot_value_array(1, predictions_single[0], test_labels)
_ = plt.xticks(range(10), class_names, rotation=45)
plt.show()



In [30]: np.argmax(predictions_single[0])

Out[30]: 2